

Week5-1: RAG — Reranking & Context Optimization

0) Setup (Colab-friendly)

```
0) Setup (Colab-friendly)
os, time, math, random, re
from typing import List, Dict, Tuple
import numpy as np, pandas as pd, matplotlib.pyplot as plt
from collections import Counter
DATA_DIR = './week5_data'; os.makedirs(DATA_DIR, exist_ok=True)
CORPUS_CSV = os.path.join(DATA_DIR, '/content/sample_medical_texts.csv')
if not os.path.exists(CORPUS_CSV):
    demo = pd.DataFrame({
        'doc_id': [f'doc{i}' for i in range(1, 11)],
        'chunk_id': [f'doc{i}_chunk1' for i in range(1, 11)],
        'text': [f'This is demo text about topic {i}. It mentions entity E{i%3} and method M{i%2}.' for i in range(1, 11)]
    })
df_demo.to_csv(CORPUS_CSV, index=False)
corpus = pd.read_csv(CORPUS_CSV); print('✅ Corpus loaded:', corpus.shape); display(corpus.head(3))
```

✅ Corpus loaded: (5, 5)

	doc_id	original_text	summary	validated_summary	rating
0	1	Patient reports chest pain radiating to the le...	NaN	NaN	NaN
1	2	The subject is a 45-year-old male with a histo...	NaN	NaN	NaN
2	3	MRI scan reveals a mass in the right temporal ...	NaN	NaN	NaN

```
# 1) Baseline + Dense + RRF
def simple_tokenize(text):
    return re.findall(r"[A-Za-z0-9_]+", text.lower())
docs = corpus['original_text'].tolist()
doc_tokens = [simple_tokenize(t) for t in docs]
from collections import Counter
doc_counts = [Counter(t) for t in doc_tokens]
N = len(docs); import math
avgdl = sum(len(t) for t in doc_tokens)/max(N,1)
def bm25_score(query, idx, k1=1.5, b=0.75):
    q_tokens = simple_tokenize(query); score = 0.0; dl = len(doc_tokens[idx])
    for w in q_tokens:
        n_w = sum(1 for c in doc_tokens if w in c)
        if n_w == 0: continue
        idf = math.log((N - n_w + 0.5)/(n_w + 0.5) + 1)
        f = doc_counts[idx][w]; denom = f + k1*(1 - b + b*dl/(avgdl+1e-9))
        score += idf * (f*(k1+1))/(denom + 1e-9)
    return score
import numpy as np
np.random.seed(7); dense_dim = 128
dense_index = np.random.randn(N, dense_dim).astype('float32')
dense_index /= (np.linalg.norm(dense_index, axis=1, keepdims=True) + 1e-9)
def bm25_search(query, k=5):
    scores = [(i, bm25_score(query, i)) for i in range(N)]; scores.sort(key=lambda x: -x[1]); return scores[:k]
def dense_search(query, k=5):
    q_vec = np.random.randn(dense_dim).astype('float32'); q_vec /= (np.linalg.norm(q_vec)+1e-9)
    sims = dense_index @ q_vec; top = np.argsort(-sims)[:k];
    return [(int(i), float(sims[i])) for i in top]
def rrf_fuse(bm_hits, de_hits, k=60):
    rank = {}
    for r, (i, _) in enumerate(bm_hits): rank[i] = rank.get(i,0) + 1.0/(k + r + 1)
    for r, (i, _) in enumerate(de_hits): rank[i] = rank.get(i,0) + 1.0/(k + r + 1)
    return sorted(rank.items(), key=lambda x: -x[1])
print('BM25 sample:', bm25_search('topic 1',5))
print('Dense sample:', dense_search('topic 1',5))
print('RRF sample:', rrf_fuse(bm25_search('topic 1',5), dense_search('topic 1',5))[:5])
```

```
BM25 sample: [(0, 0.0), (1, 0.0), (2, 0.0), (3, 0.0), (4, 0.0)]
Dense sample: [(4, 0.11497273296117783), (1, 0.05473370850086212), (2, 0.041441410779953), (3, 0.017484545707702637),
RRF sample: [(0, 0.03252247488101534), (2, 0.032266458495966696), (1, 0.0315136476426799), (3, 0.03149801587301587), (
```

```
# 2) Rerank + MMR + compression
def rerank(query, candidates, top_k=5):
    q = np.random.randn(dense_dim).astype('float32')
    q /= (np.linalg.norm(q)+1e-9)
```

```

scored = []
for (idx, _) in candidates:
    s = float(dense_index[idx] @ q + 0.05*np.random.randn()); scored.append((idx, s))
scored.sort(key=lambda x: -x[1]); return scored[:top_k]
def mmr_select(query, candidates, lambda_mult=0.6, k=4):
    if not candidates: return []
    q = np.random.randn(dense_dim).astype('float32'); q /= (np.linalg.norm(q)+1e-9)
    cand_vecs = dense_index[[idx for idx in candidates]]
    selected, remaining = [], list(range(len(candidates)))
    while remaining and len(selected) < k:
        best, best_score = None, -1e9
        for j in remaining:
            rel = float(cand_vecs[j] @ q)
            div = 0.0 if not selected else max(float(cand_vecs[j] @ cand_vecs[s]) for s in selected)
            score = lambda_mult*rel - (1-lambda_mult)*div
            if score > best_score: best, best_score = j, score
        selected.append(best); remaining.remove(best)
    return [candidates[j] for j in selected]
def compress_passage(text, max_tokens=60):
    toks = text.split(); return ' '.join(toks[:max_tokens])
query='Explain topic 1 and method M0'
bm = bm25_search(query, k=8); de = dense_search(query, k=8)
fused = [i for i, _ in rrf_fuse(bm, de)[:8]]
reranked = [i for i, _ in rerank(query, [(i,0) for i in fused], top_k=5)]
mmr = mmr_select(query, fused, lambda_mult=0.6, k=4)
print('Fused:', fused[:8]); print('Reranked:', reranked); print('MMR:', mmr)

```

```

Fused: [3, 0, 4, 1, 2]
Reranked: [0, 2, 4, 3, 1]
MMR: [3, 0, 2, 1]

```

```

import time
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

# -----
# Example Query Set
# -----
QUERIES = [
    "Explain how Reciprocal Rank Fusion improves retrieval performance.",
    "What is the role of the AgentManager entity?",
    "Summarize the reproducibility practices in Week 3 experiments."
]

# -----
# Dummy pipeline + recall function
# (Replace with your actual implementations)
# -----
def pipeline(query, variant="baseline", top_k=5, compress=False):
    """
    Simulates retrieval pipeline.
    Returns (idxs, contexts).
    """
    idxs = list(range(top_k))
    contexts = [f"context for {query} (variant={variant}, compress={compress})" for _ in idxs]
    return idxs, contexts

def recall_like(query, idxs):
    """
    Dummy recall-like metric.
    """
    return round(len(idxs) / 10, 2)

# -----
# Safe validator agent
# -----
class SummarizeValidatorAgent:
    def __init__(self, verbose=True):
        self.verbose = verbose

    def _log(self, msg):
        if self.verbose:
            print(f"[SummarizeValidatorAgent] {msg}")

    def validate_summary(self, original_text: str, summary: str):
        if not original_text or not summary:

```

```

        return {"analysis": "Missing input text or summary", "rating": 1}

    analysis = (
        f"Original length: {len(original_text)} chars, "
        f"Summary length: {len(summary)} chars."
    )
    rating = 5 if len(summary) < len(original_text) else 3
    return {"analysis": analysis, "rating": rating}

def execute(self, df: pd.DataFrame):
    # Normalize column names
    df.columns = df.columns.str.lower().str.strip()

    # Rename if common variations found
    rename_map = {}
    if "text" in df.columns:
        rename_map["text"] = "original_text"
    if "input" in df.columns:
        rename_map["input"] = "original_text"
    if "summarized_text" in df.columns:
        rename_map["summarized_text"] = "summary"
    if "output" in df.columns:
        rename_map["output"] = "summary"
    if rename_map:
        df = df.rename(columns=rename_map)

    # Check required columns
    if not {"original_text", "summary"}.issubset(df.columns):
        raise KeyError("DataFrame must contain 'original_text' and 'summary' columns.")

    results = []
    for _, row in df.iterrows():
        orig = row["original_text"]
        summ = row["summary"]
        res = self.validate_summary(orig, summ)
        results.append(res)

    return pd.DataFrame(results)

# -----
# Main Evaluation Loop
# -----
rows = []
for q in QUERIES:
    for variant, comp in [
        ("baseline", False),
        ("rerank", False),
        ("compress", True),
        ("rerank+compress", True)
    ]:
        t0 = time.time()
        idxs, ctx = pipeline(q, variant=variant, top_k=5, compress=comp)
        t1 = time.time()

        rows.append({
            "query": q,
            "variant": variant,
            "recall_like": recall_like(q, idxs),
            "latency_s": round(t1 - t0, 4),
            "avg_ctx_len": round(np.mean([len(c) for c in ctx]), 1)
        })

results_df = pd.DataFrame(rows)
print("\n=== Pipeline Evaluation Results ===")
print(results_df)

# -----
# Example usage of validator
# -----
data = {
    "text": [
        "Machine learning models are widely used for classification tasks.",
        "Blockchain can enhance transparency in supply chains."
    ],
    "summarized_text": [
        "ML models help in classification.",
        "Blockchain improves supply chain transparency."
    ]
}

```

```

    ]
}
df = pd.DataFrame(data)

agent = SummarizeValidatorAgent(verbose=True)
validated = agent.execute(df)
print("\n=== Validation Results ===")
print(validated)

=== Pipeline Evaluation Results ===

```

	query	variant \
0	Explain how Reciprocal Rank Fusion improves re...	baseline
1	Explain how Reciprocal Rank Fusion improves re...	rerank
2	Explain how Reciprocal Rank Fusion improves re...	compress
3	Explain how Reciprocal Rank Fusion improves re...	rerank+compress
4	What is the role of the AgentManager entity?	baseline
5	What is the role of the AgentManager entity?	rerank
6	What is the role of the AgentManager entity?	compress
7	What is the role of the AgentManager entity?	rerank+compress
8	Summarize the reproducibility practices in Wee...	baseline
9	Summarize the reproducibility practices in Wee...	rerank
10	Summarize the reproducibility practices in Wee...	compress
11	Summarize the reproducibility practices in Wee...	rerank+compress

	recall_like	latency_s	avg_ctx_len
0	0.5	0.0	113.0
1	0.5	0.0	111.0
2	0.5	0.0	112.0
3	0.5	0.0	119.0
4	0.5	0.0	91.0
5	0.5	0.0	89.0
6	0.5	0.0	90.0
7	0.5	0.0	97.0
8	0.5	0.0	109.0
9	0.5	0.0	107.0
10	0.5	0.0	108.0
11	0.5	0.0	115.0


```

=== Validation Results ===

```

	analysis	rating
0	Original length: 65 chars, Summary length: 33 ...	5
1	Original length: 53 chars, Summary length: 46 ...	5

```

import matplotlib.pyplot as plt


# -----
# Plot Recall-like metric
# -----
plt.figure(figsize=(10, 6))
results_df.groupby("variant")["recall_like"].mean().plot(
    kind="bar", color="skyblue", edgecolor="black"
)
plt.title("Average Recall-like Score by Variant")
plt.ylabel("Recall-like")
plt.xlabel("Variant")
plt.xticks(rotation=30)
plt.grid(axis="y", linestyle="--", alpha=0.7)
plt.show()

# -----
# Plot Latency
# -----
plt.figure(figsize=(10, 6))
results_df.groupby("variant")["latency_s"].mean().plot(
    kind="bar", color="lightgreen", edgecolor="black"
)
plt.title("Average Latency by Variant")
plt.ylabel("Seconds")
plt.xlabel("Variant")
plt.xticks(rotation=30)
plt.grid(axis="y", linestyle="--", alpha=0.7)
plt.show()

# -----
# Plot Avg Context Length
# -----
plt.figure(figsize=(10, 6))
results_df.groupby("variant")["avg_ctx_len"].mean().plot(

```

```
        kind="bar", color="salmon", edgecolor="black"  
    )  
plt.title("Average Context Length by Variant")  
plt.ylabel("Characters")  
plt.xlabel("Variant")  
plt.xticks(rotation=30)  
plt.grid(axis="y", linestyle="--", alpha=0.7)  
plt.show()
```

```
print("Mini-Report: fill in observations and best recipe.")
```

```
Mini-Report: fill in observations and best recipe.
```

```
%pip install -q -U langchain_community langchain-google-genai pypdf
```

```
import os, pathlib
from langchain_community.vectorstores import Chroma
from langchain_community.document_loaders import PyPDFLoader
from langchain.text_splitter import RecursiveCharacterTextSplitter
from langchain_google_genai import GoogleGenerativeAIEmbeddings
```

```
import getpass
```

```
if not os.getenv("GEMINI_API_KEY"):
    os.environ["GEMINI_API_KEY"] = getpass.getpass("AIzaSyAgmfCXklGAKx4xXFwmfefKhIjtdpDSq4w ")
```

```
AIzaSyAgmfCXklGAKx4xXFwmfefKhIjtdpDSq4w .....
0.02
```

```
pdf_path = "/content/Week4-RAG_LangChain_AppliedGenAI.ipynb - Colab.pdf" # 📄 replace with your Week 4 file
loader = PyPDFLoader(pdf_path)
docs = loader.load()
```

```
print("✅ Loaded", len(docs), "PDF pages")
```

```
✅ Loaded 41 PDF pages
```

```
splitter = RecursiveCharacterTextSplitter(
    chunk_size=500, chunk_overlap=100
)
chunks = splitter.split_documents(docs)

print("✅ Split into", len(chunks), "chunks")
```

```
✅ Split into 41 chunks
```

```
%pip install -q -U google-generativeai langchain-google-genai
```

```
1.3/1.3 MB 11.2 MB/s eta 0:00:00
```

```
import os, getpass
import google.generativeai as genai
```

```
# Securely input your API key (from https://aistudio.google.com/app/apikey)
if not os.getenv("GEMINI_API_KEY"):
    os.environ["GEMINI_API_KEY"] = getpass.getpass("AIzaSyBdFwoE0bkuNWcc5aMMh50BCw1CD50oC54")
```

